

HETKUMAR PATEL

Applied AI & Agentic Systems

Toronto, ON (Eastern Time) (647) 835-3408 hetkumar.patell1403@gmail.com

linkedin.com/in/h3t08 github.com/het0814 het.codes

Summary

Applied AI software engineer who designs, builds, and ships production-grade, LLM-powered agentic systems from prototype to deployment. Hands-on across the full LLM stack, frontier models (Claude, GPT, Gemini), RAG, multi-step planning agents (ReAct / Plan-and-Execute), MCP and multi-agent orchestration with a strong focus on rigorous evaluation frameworks that measure agent accuracy, safety and latency beyond trial-and-error. Writes clean, testable, observable Python and flexes into frontend (Next.js / React) when the problem demands it. Co-author of three AI/ML papers; comfortable working directly with stakeholders to translate ambiguous business problems into reliable, auditable solutions, including in regulated domains like healthcare.

Experience

AI Engineer

Jan. 2026 – Present

Wedge (Y Combinator S25)

San Francisco, CA · Remote

- Led production LLM agents for healthcare workflows deployed with 3 clinical providers on live patient data (250–300 referrals/mo), directing 2 interns to automate referral intake, follow-up and payer communication in a regulated, HIPAA-sensitive environment.
- Shipped an Automated Referral System end-to-end that cut manual work by 90% and saved 20+ hrs/week, shifting staff from hands-on processing to lightweight monitoring: fax ingestion + LLM-assisted parsing (pdfplumber + Groq), structured records in Supabase, and a triage agent + UI, with Playwright automation for payer portals.
- Built an outbound voice agent on LiveKit (Deepgram STT + ElevenLabs TTS + OpenAI reasoning) that autonomously calls insurance payers for verification and prior-authorization follow-up, tuned for low-latency turn-taking on live calls.
- Implemented Twilio + APScheduler SMS/voice follow-up and a knowledge layer feeding the agents, in clean Python (FastAPI, Pydantic).

AI Software Developer

Feb. 2026 – Present

Moorcheh.ai by Edge AI Innovations

Toronto, ON · Remote

- Built and shipped Memanto, an AI memory agent, end-to-end on Moorcheh's architecture (18K PySDK downloads); co-authored an arXiv paper on its typed semantic memory and information-theoretic retrieval for long-horizon agents.
- Maintain the open-source repo (1.8K stars, 604 forks, 200+ PRs), engineering the core memory pipeline on RAG semantic search to persist and retrieve context across sessions, exposed as a first-class tool via MCP for multi-agent workflows.
- Iterated on retrieval, latency and reliability for a production agentic system, reaching 87.1 on LoCoMo and 89.8 on LongEval; shipped the Next.js landing + MDX docs sites.

Gen AI Associate: LLM Evaluation

Jan. 2026 – Mar. 2026

Innodata Inc.

Toronto, ON · Remote

- Ran structured LLM evaluation workflows across locality, freshness, completeness and on-point accuracy, applying A/B-style testing and human-in-the-loop review to measure and improve model output quality.
- Contributed evaluation insights to large-scale model-improvement pipelines across multiple model versions.

Artificial Intelligence Researcher (CO-OP, Part-Time)

Sep. 2024 – Aug. 2025

Sheridan College, Centre for Applied AI

Oakville, ON · Hybrid

- Built an end-to-end RAG pipeline on LangChain with vector indexing for KPI extraction from unstructured fleet documents, and fine-tuned LLMs (LoRA / PEFT) on domain corpora with documented evaluation benchmarks.
- Deployed agentic workflows with function calling, tool orchestration and automated quality checks; co-authored two peer-reviewed papers presented at international conferences and mentored junior researchers.

AI/ML Software Developer (CO-OP)

May 2024 – Sep. 2024

Naryant, Centre for Applied AI

Oakville, ON · Hybrid

- Owned the full ML lifecycle for a transportation mode detection system in Python, Scikit-learn, TensorFlow (92% accuracy), and built a CTGAN synthetic-data pipeline for SUMO mobility simulations.

IIoT Developer (CO-OP)

May 2023 – Feb. 2024

Magna International

Milton, ON · On-site

- Delivered real-time monitoring (Ignition + Python) that cut troubleshooting time by 85% and improved monitoring accuracy by 97% across production machinery.

Projects

ContextIQ: Production RAG System | *FastAPI, LangChain, Pinecone, pgvector, React, Docker* | **GitHub** **Mar. 2025**

- End-to-end **RAG application** with document ingestion, hybrid vector search (Pinecone + pgvector), re-ranking and query rewriting; FastAPI backend with **groundedness scoring** and **structured-output validation** for auditability, containerized with Docker.

Restro: Multilingual Voice Agent | *LiveKit, Claude, Deepgram, Cartesia, Twilio SIP* **2026**

- Real-time multilingual restaurant voice agent (English, Hindi, Gujarati, Punjabi) with full pipeline: SIP ingress, VAD, streaming STT, LLM reasoning, low-latency TTS and turn-taking logic.

Technical Skills

Agentic & LLM: Agent design (ReAct, Plan-and-Execute), MCP, multi-agent orchestration (LangChain, LangGraph, CrewAI, AutoGen, Semantic Kernel), tool/function calling, RAG, semantic memory, prompt engineering, guardrails, LoRA/PEFT fine-tuning

LLM APIs & Frameworks: Anthropic Claude, OpenAI, Google Gemini, Vertex AI, Groq; Hugging Face, PyTorch, TensorFlow, Scikit-learn, llama.cpp

Evaluation & MLOps: Evaluation frameworks (accuracy/safety/latency), A/B testing, human-in-the-loop, groundedness scoring, LongMem/LoCoMo benchmarks, pytest, MLflow, model versioning, CI/CD, GitHub Actions, Docker, Kubernetes

Backend & Frontend: Python (FastAPI, asyncio, Pydantic), Flask, REST APIs; TypeScript, Next.js, React, Tailwind; C#, ASP.NET

Vector & Data: Pinecone, pgvector, Chroma, Milvus, Redis, hybrid search, re-ranking; SQL, MongoDB, Supabase, Pandas, Spark, Databricks

Voice & Cloud: LiveKit, Deepgram, ElevenLabs, Twilio SIP, Silero VAD; AWS, Azure, GCP, Terraform, Vercel

Research, Education & Certifications

Publications: (1) Memanto: Typed Semantic Memory with Information-Theoretic Retrieval for Long-Horizon Agents (arXiv); (2) Integrating Knowledge & Data-Driven Methods to Generate Synthetic Mobility Data (IEEE); (3) Topic Modeling Enhancement using LLM-Generated Summaries (IEEE)

Education: Honours Bachelor of Computer Science (Data Analytics), Sheridan College – GPA 3.7, Awarded Honours (2025)

Certifications: Prompt Engineering & Programming with OpenAI, Machine Learning I — Columbia (2025); OCI Generative AI Professional, Data Science & AI Foundations — Oracle (2025); Generative AI Fundamentals — Databricks; Building GenAI Applications — MongoDB; AWS Educate: Intro to Generative AI; GitHub Foundations — Microsoft (2025)